

## Single Agent Architecture as Field Theory: Technical TLDR

Guillem Duran Ballester, Jan 2026

## 0. Architecture

The following diagrams illustrate the current implementation architecture of the TopoEncoder system, showing how observations are encoded into the split-latent space ( $K, z_n, z_{tex}$ ) and decoded back to reconstructions. These diagrams mirror the code in `src/fragile/core/layers/atlas.py` and `src/experiments/topoencoder_2d.py`.

## 0.1 CovariantChartRouter

The chart router is shared by both encoder and decoder. It performs metric-aware, covariant chart assignment using:

- Geodesic query terms (linear + quadratic Christoffel correction)
- Wilson-line transport via Cayley transform for gauge invariance
- Position-dependent temperature scaled by local metric (conformal factor)

```

classDef io fill:#0b1320,stroke:#93c5fd,stroke-width:1px,color:#e5e7eb;
classDef feat fill:#111827,stroke:#22d3ee,stroke-width:1px,color:#e5e7eb;
classDef router fill:#2b1f1f,stroke:#f59e0b,stroke-width:1px,color:#e5e7eb;
classDef geom fill:#1f2937,stroke:#a78bfa,stroke-width:1px,color:#e5e7eb;
classDef util fill:#262626,stroke:#a3a3a3,stroke-width:1px,color:#e5e7eb;

class Z,Kchart geom;
class F,Qfeat feat;
class Qz,Gamma,Qsum,Transport,Cayley,ChartTokens,KeyProj,ChartQ,KeyMerge,Keys,Scores,Tau,Sum util;

```

## 0.2 Full TopoEncoder (Encoder + Decoder)

Complete end-to-end architecture showing: - **Encoder**: Feature extraction → Chart routing → VQ per chart → Split into  $(z_{geo}, z_n, z_{tex})$  - **Decoder**: Chart-weighted reconstruction → Geometric base + Texture residual

```

%{init: {"themeVariables": {"background": "#0b111b", "edgeLabelBackground": "#111827", "textColor": "#e5e7eb", "fontColor": "#e5e7eb", "fontStyle": "bold", "fontSize": 14, "fontWeight": "bold"}, "fontFamily": "monospace", "fontSize": 12, "fontStyle": "normal", "fontWeight": "normal"}, "script": "js", "language": "mermaid", "id": "graph TD", "label": "Flowchart of the TopoEncoder architecture showing the flow from input features to final geometric and texture outputs."}
flowchart TD
    subgraph TOP ["TopoEncoderPrimitives (current code)"]
        subgraph ENC ["PrimitiveAttentiveAtlasEncoder"]
            X["Input x [B, D_in]"] -- "x [B, D_in]" --> FE["Feature extractor\\nMLP: SpectralLinear"]
            FE -- "features [B, H]" --> F["features [B, H]"]
            F -- "features [B, H]" --> Vproj["val_proj: SpectralLinear\\nv [B, D]"]
            ChartCenters["chart_centers c_k [N_c, D]"] -- "c_k [N_c, D]" --> RouterEnc["Chart Router"]
            RouterEnc -- "w_enc [B, N_c]" --> Wenc["w_enc [B, N_c]"]
            RouterEnc -- "K_chart [B]" --> Kchart["K_chart [B]"]

            Wenc -- "w_enc [B, N_c]" --> Cbar["c_bar = sum(w_enc * c_k) [B, D]"]
            ChartCenters -- "c_k [N_c, D]" --> Cbar
            Vproj -- "v [B, D]" --> Vlocal["v_local = v - c_bar [B, D]"]
            Cbar -- "c_bar [B, D]" --> Vlocal

            Codebook["Codebook (deltas) [N_c, K, D]"] -- "codebook [N_c, K, D]" --> Diff["diff"]
            Vlocal -- "v_local [B, D]" --> Diff
            Diff -- "diff [B, N_c, K, D]" --> SoftEq["SoftEquivariantLayer per chart\\n(optional)"]
            SoftEq -- "diff' [B, N_c, K, D]" --> Dist["dist = ||diff'||^2 [B, N_c, K]"]
            Diff -.--> Dist

            Dist -- "dist [B, N_c, K]" --> Indices["indices per chart [B, N_c]"]
            Indices -- "indices [B, N_c]" --> ZqAll["z_q_all [B, N_c, D]\\n(gather; + soft-ST i..."]
            ZqAll -- "z_q_all [B, N_c, D]" --> ZqBlend["z_q_blended = sum(w_enc * z_q_all)"]

            Indices -- "indices [B, N_c]" --> Kcode["K_code (from K_chart)"]
            Kchart -- "K_chart [B]" --> Kcode
        end
    end

```

```

ZqAll -- "z_q_all [B, N_c, D]" --> VQLoss["vq_loss = codebook + 0.25 * commitment"]
Vlocal -- "v_local [B, D]" --> VQLoss

ZqAll -- "z_q_all [B, N_c, D]" --> DeltaAll["delta_all = v_local - z_q_all (detach)
DeltaAll -- "delta_all [B, N_c, D]" --> Struct["structure_filter\nIsotropicBlock +
Struct -- "z_n_all_charts [B, N_c, D]" --> ZnAll["z_n_all_charts [B, N_c, D]"]
ZnAll -- "z_n_all_charts [B, N_c, D]" --> Zn["z_n = sum(w_enc * z_n_all_charts) [B
ZqBlend -- "z_q_blended [B, D]" --> DeltaBlend["delta_blended = v_local - z_q_blend
DeltaBlend -- "delta_blended [B, D]" --> Ztex["z_tex = delta_blended - z_n"]

ZqBlend -- "z_q_blended [B, D]" --> ZqSt["z_q_st = v_local + (z_q_blended - v_local
ZqSt -- "z_q_st [B, D]" --> Zgeo["z_geo = c_bar + z_q_st + z_n"]
Zn -- "z_n [B, D]" --> Zgeo
Cbar -- "c_bar [B, D]" --> Zgeo

ZnAll -- "z_n_all_charts [B, N_c, D]" --> Jump["FactorizedJumpOperator (optional)"]
end

subgraph DEC["PrimitiveTopologicalDecoder"]
Zgeo -- "z_geo [B, D]" --> TanhG["tanh(z_geo)"]
TanhG -- "tanh(z_geo) [B, D]" --> RouterDec["Chart router\nCovariantChartRouter (co
RouterDec -- "w_dec [B, N_c]" --> Wdec["w_dec [B, N_c]"]
ChartIdx["chart_index (optional)"] -- "K_chart [B]" --> OneHot["one-hot -> w_hard"]
OneHot -- "w_dec_hard [B, N_c]" --> Wdec

TanhG -- "tanh(z_geo) [B, D]" --> ChartProj["chart_projectors: SpectralLinear x N_c
ChartProj -- "h_i [B, N_c, H]" --> Gate["NormGatedGELU on h_stack"]
Gate -- "h_stack [B, N_c, H]" --> Mix["h_global = sum(w_dec * h_stack)"]
Wdec -- "w_dec [B, N_c]" --> Mix

Mix -- "h_global [B, H]" --> Renderer["renderer: SpectralLinear + NormGatedGELU x2
Mix -- "h_global [B, H]" --> Skip["render_skip: SpectralLinear"]
Renderer -- "h_render [B, D_out]" --> AddSkip["x_hat_base = renderer + skip"]
Skip -- "h_skip [B, D_out]" --> AddSkip

Ztex -- "z_tex [B, D]" --> TanhT["tanh(z_tex)"]
TanhT -- "tanh(z_tex) [B, D]" --> TexRes["tex_residual: SpectralLinear"]
TexRes -- "tex_resid [B, D_out]" --> AddTex["x_hat = x_hat_base + tex_residual_sca
AddSkip -- "x_hat_base [B, D_out]" --> AddTex
AddTex -- "x_hat [B, D_out]" --> Xhat["x_hat [B, D_out]"]
end
end

classDef io fill:#0b1320,stroke:#93c5fd,stroke-width:1px,color:#e5e7eb;
classDef feat fill:#111827,stroke:#22d3ee,stroke-width:1px,color:#e5e7eb;
classDef router fill:#2b1f1f,stroke:#f59e0b,stroke-width:1px,color:#e5e7eb;
classDef vq fill:#1f2f2a,stroke:#34d399,stroke-width:1px,color:#e5e7eb;
classDef geom fill:#1f2937,stroke:#a78bfa,stroke-width:1px,color:#e5e7eb;

```

```
classDef residual fill:#3b1f2b,stroke:#f472b6,stroke-width:1px,color:#e5e7eb;
classDef decoder fill:#1f2b3b,stroke:#60a5fa,stroke-width:1px,color:#e5e7eb;
classDef util fill:#262626,stroke:#a3a3a3,stroke-width:1px,color:#e5e7eb;

class X,Xhat,ChartIdx,Kchart io;
class FE,F,Vproj,Struct feat;
class RouterEnc,RouterDec,Wenc,Wdec,OneHot router;
classs Codebook,Diff,SoftEq,Dist,Indices,ZqAll,ZqBlend,VQLoss,Kcode vq;
class ChartCenters,Cbar,Vlocal,Zgeo,ZqSt,ZnAll,Zn geom;
class DeltaAll,DeltaBlend,Ztex,TanhT,TexRes residual;
class TanhG,ChartProj,Gate,Mix,Renderer,Skip,AddSkip,AddTex decoder;
class Jump util;
```

### 0.3 Decoder Detail (Inverse Atlas)

Focused view of the decoder showing how the geometric latent  $z_{geo}$  and texture residual  $z_{tex}$  are independently processed and combined:

- Geometric path: Chart projectors → Chart-weighted mixing → Renderer
- Texture path: Independent residual network
- Final output: Base reconstruction + scaled texture residual

```

%%{init: {"themeVariables": {"background": "#0b111b", "edgeLabelBackground": "#111827", "textColor": "#e5e7eb"}, "graph": {"fontFamily": "monospace", "fontSize": 14, "nodeSize": 1000, "edgeWidth": 2}, "node": {"shape": "rect", "strokeWidth": 1, "stroke": "#93c5fd", "fill": "#0b111b", "color": "#e5e7eb"}, "edge": {"strokeWidth": 1, "stroke": "#f59e0b", "color": "#e5e7eb"}, "label": {"color": "#e5e7eb"}, "highlight": {"strokeWidth": 2, "stroke": "#f59e0b", "fill": "#0b111b", "color": "#e5e7eb"}, "highlightLabel": {"color": "#e5e7eb"}}, "graph TD", "graph LR"]}

flowchart TD
    subgraph DEC["PrimitiveTopologicalDecoder (current code)"]
        Zgeo["z_geo = c_bar + z_q_st + z_n [B, D]"] -- "z_geo [B, D]" --> TanhG["tanh(z_geo)"]
        TanhG -- "tanh(z_geo) [B, D]" --> RouterDec["Chart router\nCovariantChartRouter (covariant chart router)"]
        RouterDec -- "w_dec [B, N_c]" --> Wdec["w_dec [B, N_c]"]
        ChartIdx["chart_index (optional)"] -- "K_chart [B]" --> OneHot["one-hot -> w_hard"]
        OneHot -- "w_dec_hard [B, N_c]" --> Wdec

        TanhG -- "tanh(z_geo) [B, D]" --> ChartProj["chart_projectors: SpectralLinear x N_c"]
        ChartProj -- "h_i [B, N_c, H]" --> Gate["NormGatedGELU on h_stack"]
        Gate -- "h_stack [B, N_c, H]" --> Mix["h_global = sum(w_dec * h_stack)"]
        Wdec -- "w_dec [B, N_c]" --> Mix

        Mix -- "h_global [B, H]" --> Renderer["renderer: SpectralLinear + NormGatedGELU x2 + SpatialConvolution"]
        Mix -- "h_global [B, H]" --> Skip["render_skip: SpectralLinear"]
        Renderer -- "h_render [B, D_out]" --> AddSkip["x_hat_base = renderer + skip"]
        Skip -- "h_skip [B, D_out]" --> AddSkip

        Ztex["z_tex [B, D]"] -- "z_tex [B, D]" --> TanhT["tanh(z_tex)"]
        TanhT -- "tanh(z_tex) [B, D]" --> TexRes["tex_residual: SpectralLinear"]
        TexRes -- "tex_resid [B, D_out]" --> AddTex["x_hat = x_hat_base + tex_residual_scale * tex_resid"]
        AddSkip -- "x_hat_base [B, D_out]" --> AddTex
        AddTex -- "x_hat [B, D_out]" --> Xhat["x_hat [B, D_out]"]
    end

    classDef io fill:#0b1320,stroke:#93c5fd,stroke-width:1px,color:#e5e7eb;
    classDef router fill:#2b1f1f,stroke:#f59e0b,stroke-width:1px,color:#e5e7eb;

```

```

classDef geom fill:#1f2937,stroke:#a78bfa,stroke-width:1px,color:#e5e7eb;
classDef residual fill:#3b1f2b,stroke:#f472b6,stroke-width:1px,color:#e5e7eb;
classDef decoder fill:#1f2b3b,stroke:#60a5fa,stroke-width:1px,color:#e5e7eb;

class Zgeo geom;
class RouterDec,Wdec,OneHot router;
class ChartIdx,Xhat io;
class TanhG,ChartProj,Gate,Mix,Renderer,Skip,AddSkip,AddTex decoder;
class Ztex,TanhT,TexRes residual;

```

#### 0.4 Experiment Wiring (Training Losses)

Optional training components available in the experiment configuration:

- Learned precisions for automatic loss balancing (reconstruction, VQ, supervised)
- SupervisedTopologyLoss for semantic topology learning
- FactorizedJumpOperator for chart transition consistency
- InvariantChartClassifier (detached readout head with separate optimizer)

```

%%{init: {"themeVariables": {"background": "#0b111b", "edgeLabelBackground": "#111827", "textColor": "#e5e7eb"}}

graph TD
    X["batch_X"] --> Enc["TopoEncoderPrimitives.encoder"]
    Enc -- "z_geo [B, D]" --> Dec["TopoEncoderPrimitives.decoder"]
    Enc -- "z_tex [B, D]" --> Dec
    Dec -- "recon_a [B, D_in]" --> ReconLoss["recon_loss = mse(recon_a, batch_X)"]
    X --> ReconLoss

    Enc -- "vq_loss" --> VQLoss["vq_loss (codebook + commitment)"]

    ReconLoss --> ReconTerm["recon_term\n(optional learned precision)"]
    VQLoss --> VQTerm["vq_term\n(optional learned precision)"]

    Enc -- "enc_w [B, N_c]" --> Sup["SupervisedTopologyLoss (optional)"]
    Enc -- "z_geo [B, D]" --> Sup
    Y["batch_labels [B]"] --> Sup
    Sup -- "sup_total + components" --> SupTerm["sup_term\n(optional learned precision)"]

    Enc -- "z_n_all_charts [B, N_c, D]" --> Jump["FactorizedJumpOperator (optional)"]
    Enc -- "enc_w [B, N_c]" --> Jump
    Jump -- "jump_loss (schedule weight)" --> LossA["atlas loss\n(recon + vq + regs + jump + sup)"]

    ReconTerm --> LossA
    VQTerm --> LossA
    SupTerm -- "sup_weight * sup_term" --> LossA

    Enc -- "enc_w (detach)" --> Cls["InvariantChartClassifier (optional)"]
    Enc -- "z_geo (detach)" --> Cls
    Y --> CE["cross_entropy"]
    Cls -- "logits [B, C]" --> CE
    CE --> OptCls["opt_classifier.step()"]

```

## 1. Latent Space Decomposition

**Split-Latent Structure:**

- Total latent:  $Z = (K, z_n, z_{\text{tex}})$  where each component has distinct geometric/physical role
- $K \in \{1, \dots, N_c\}$ : Discrete macro state (VQ codebook index) - chart assignment on topological atlas
- $z_n \in \mathbb{R}^{D_n}$ : Continuous nuisance latent - local coordinates within chart (gauge-invariant position)
- $z_{\text{tex}} \in \mathbb{R}^{D_t}$ : Texture residual - holographic boundary degrees of freedom
- Decomposition satisfies  $z_e = z_q + z_n + z_{\text{tex}}$  where  $z_e$  is raw encoder output,  $z_q$  is VQ-quantized macro

**Encoder Architecture (AttentiveAtlasEncoder):**

- Feature extraction: Conv layers ( $64 \rightarrow 128 \rightarrow 256$  channels)  $\rightarrow$  hidden\_dim projection
- Cross-attention routing:  $w_k = \text{softmax}(\langle K_k, Q(x) \rangle / \sqrt{D})$  where  $K_k$  are learnable chart queries
- Query projection:  $Q(x) = \text{key\_proj}(\text{features}(x))$  with LayerNorm stabilization
- Chart assignment:  $K = \arg \max_k w_k(x)$
- VQ per chart:  $N_c$  independent codebooks, each with  $K_c$  codes, vectorized quantization
- Nuisance extraction: Structure filter  $z_n = f_{\text{struct}}(z_e - z_q)$  removes VQ-residual structure
- Texture: Holographic residual  $z_{\text{tex}} = (z_e - z_q) - z_n$  orthogonal to both macro and nuisance.
- **Texture Firewall:** Dynamics depend on  $z_{\text{macro}}$  and  $z_n$ , screening out only  $z_{\text{tex}}$ .

**Training Objectives for Coarse-Graining:** To enforce this split-latent structure and ensure valid coarse-graining, we minimize:

$$\mathcal{L}_{\text{latent}} = \mathcal{L}_{\text{VQ}} + \mathcal{L}_{\text{closure}} + \mathcal{L}_{\text{slowness}} + \mathcal{L}_{\text{disentangle}}$$

1. **VQ Loss:**  $\|z_e - z_q\|^2 + \beta \|z_q - \text{sg}[z_e]\|^2$ . Stabilizes the discrete macro state  $K$ .
2. **Causal Enclosure:**  $-\log p(K_{t+1}|K_t, a_t)$ . Ensures macro dynamics are self-contained (predictable without micro details).
3. **Slowness (Anti-Churn):**  $\|e_{K_t} - e_{K_{t-1}}\|_G^2$ . Penalizes rapid flickering of the macro state to ensure temporal stability.
4. **Disentanglement:**
  - *Nuisance KL:*  $D_{\text{KL}}(q(z_n|x)\|\mathcal{N}(0, I))$ . Minimal sufficient nuisance.
  - *Texture KL:*  $D_{\text{KL}}(q(z_{\text{tex}}|x)\|\mathcal{N}(0, I))$ . Texture should contain no macro info.

**Geometric Regularization (Quality Control):** To ensure the latent space is well-conditioned (high-fidelity, isometric charts), we add:

$$\mathcal{L}_{\text{reg}} = \mathcal{L}_{\text{VICReg}} + \mathcal{L}_{\text{ortho}}$$

5. **VICReg (Self-Supervision):** Prevents collapse without negative pairs.
  - *Invariance:*  $\|z - z'\|_G^2$ . Robustness to view augmentation.
  - *Variance:*  $\max(0, \gamma - \sqrt{\text{Var}(z)})$ . Forces code utilization.
  - *Covariance:*  $C(z) \approx I$ . Decorrelates latent dimensions (whitening).
6. **Orthogonality (Chart Isometry):**  $\|W^T W - I\|_F^2$  on encoder weights. Ensures the mapping from observation to latent space is locally isometric (preserves distances), crucial for meaningful geodesic calculations.

## 2. The Reward Field & Hodge Decomposition

**Physical Context:** We model the agent as a particle with **Position and Momentum** performing a **Geodesic Random Walk** on the latent manifold. Because utility is harvested via trajectory traversal, the Reward is naturally defined as a differential **1-form** coupled to velocity, rather than a static scalar field.

**Constraint:** Reward is not a scalar  $r(z)$ , but a **1-form**  $\mathcal{R}$  that depends on direction ( $\mathcal{R}_i \dot{z}^i$ ). This requires a field-theoretic treatment of value.

**Hodge Decomposition:** The reward 1-form uniquely decomposes into three orthogonal components:

$$\mathcal{R} = \underbrace{d\Phi}_{\text{Gradient}} + \underbrace{\delta\Psi}_{\text{Solenoidal}} + \underbrace{\eta}_{\text{Harmonic}}$$

- **Gradient** ( $\Phi$ ): Optimizable scalar potential (Conservative Value).
- **Solenoidal** ( $\Psi$ ): Vector potential generating the **Value Curl** (Magnetic Field)  $\mathcal{F} = d\mathcal{R} = d\delta\Psi$ .
  - *Physics:* Just as a magnetic field  $B$  exerts a **Lorentz Force**  $v \times B$ , the Value Curl  $\mathcal{F}$  exerts a velocity-dependent force  $\mathcal{F}_{ij} \dot{z}^j$  that steers the agent to **orbit** value cycles rather than converge.
- **Harmonic** ( $\eta$ ): Topological cycles (manifold holes).

**The Screened Poisson Equation (Conservative Case):** When the **Value Curl** vanishes ( $\mathcal{F} = 0$ ), the scalar value function  $V(z)$  satisfies the **Helmholtz Equation** on the manifold:

$$(-\Delta_G + \kappa^2)V(z) = \rho_r(z)$$

where  $\Delta_G$  is the Laplace-Beltrami operator,  $\kappa$  is the screening mass, and  $\rho_r$  is the reward source density.

**The Composite Navigation Potential (Runtime):** The agent navigates a **Composite Potential**  $\Phi_{\text{eff}}$  constructed at runtime by summing learned and intrinsic signals:

$$\Phi_{\text{eff}}(z) = \underbrace{V(z)}_{\text{Learned}} + \underbrace{U(z)}_{\text{Intrinsic}} + \underbrace{\mathcal{B}_{\text{safety}}(z)}_{\text{Fixed}}$$

1. **Control** ( $V$ ): **Learned** scalar value. Drives the agent towards high-reward regions.
  - *Loss:* Helmholtz Residual (see below).
2. **Generation** ( $U$ ): **Intrinsic** entropy potential (e.g., Hyperbolic expansion  $-\log \text{vol}(z)$ ). Drives exploration away from the origin.
  - *Loss:* None (Fixed geometric prior).
3. **Safety** ( $\mathcal{B}_{\text{safety}}$ ): **Fixed** safety barrier. Hard constraints (e.g., capacity limits) modeled as high-energy walls.
  - *Loss:* None (Fixed constraint).

**Neural Hodge Decomposition (Implementation):** We approximate the Hodge components using a **Multi-Head Network** sharing a common feature backbone:

1. **Scalar Head** ( $\Phi$ ): Outputs scalar  $V(z)$ .
  - *Loss:* Helmholtz Residual on the symmetric part of the reward.
  - $\mathcal{L}_\Phi = \|V(z) - (r_{\text{sym}} + \gamma V(z'))\|^2$ .
2. **Solenoidal Head** ( $\Psi$ ): Outputs vector potential  $A(z) \in \mathbb{R}^D$  (where  $\mathcal{F} = dA$ ).

- *Loss*: Residual reconstruction on the antisymmetric part.
  - $\mathcal{L}_\Psi = \|\langle \mathcal{R}, v \rangle - (\langle \nabla \Phi, v \rangle + \langle \nabla \times A, v \rangle)\|^2$ . The curl absorbs the non-integrable reward residual.
3. **Harmonic Head ( $\eta$ )**: A set of **learnable constant 1-forms**  $\eta_k$  per chart  $k$ .
- *Mechanism*: Captures global topological currents (net flux through manifold holes) that are locally constant.
  - *Loss*: Projected residual after removing Gradient and Curl components.

**Value Function Objectives:** To define the value field  $V(z)$  as the solution to the Screened Poisson Equation, we minimize:

$$\mathcal{L}_{\text{critic}} = \underbrace{\|V(z) - (r + \gamma V(z'))\|^2}_{\text{Helmholtz Residual (TD)}} + \lambda_{\text{geo}} \underbrace{\|\nabla_G V\|^2}_{\text{Smoothness}}$$

1. **Helmholtz Residual**: Enforces the PDE source term (approx. Bellman error).
2. **Geometric Regularization**: Enforces field smoothness with respect to the manifold metric ( $\|\nabla_G V\|^2 = G^{ij} \partial_i V \partial_j V$ ).

### 3. The Policy Network (Latent Action)

The Policy acts as an **External Force Field**  $u_\pi(z)$  pushing the agent through the latent manifold. It operates in a **Latent Action Space** (the Tangent Bundle  $T\mathcal{Z}$ ), decoupling low-level motor commands from high-level intent.

#### A. The Policy Model ( $\pi_\phi$ ):

- **Role**: Symmetry-breaking control field. Converts potential energy into kinetic motion.
- **Input**: Latent State  $z_t \in \mathcal{Z}$  (Position).
- **Output**: Latent Force/Action  $u_t \in T_{z_t} \mathcal{Z}$  (Tangent Vector).
  - *Note*: This latent force is subsequently decoded into boundary motor torques action  $a_t$  by the Motor/Action Decoder (Neumann Condition).
- **Latent Action Space**: The Tangent Bundle  $T\mathcal{Z}$ . Actions are vectors “pushing” the state along geodesics.

#### B. Training Losses (Tier 1):

1. **Task Loss ( $\mathcal{L}_{\text{task}}$ )**: Standard Policy Gradient / Reinforce objective to maximize expected returns.
2. **Entropy Bonus (Ergodicity)**:  $\mathcal{L}_{\text{ent}} = -H(\pi)$ . Penalizes low entropy distributions to prevent premature mode collapse and ensure thermodynamic equilibrium.
3. **Zeno Penalty (Temporal Smoothness)**:  $\mathcal{L}_{\text{zeno}} = D_{\text{KL}}(\pi_t \parallel \pi_{t-1})$ . Penalizes infinite-frequency oscillations (Zeno behavior) to ensure physically realizable trajectories.

### 4. The World Model (Covariant Integrator)

We define the World Model not as a generic RNN, but as a **Neural Integrator** that approximates the **Lorentz-Langevin SDE**:

$$dz^k = \underbrace{(-G^{kj} \partial_j \Phi + u_\pi^k)}_{\text{Gradient + Policy}} ds + \underbrace{\beta G^{km} \mathcal{F}_{mj} \dot{z}^j ds}_{\text{Lorentz Force}} - \underbrace{\Gamma_{ij}^k \dot{z}^i \dot{z}^j ds}_{\text{Geodesic Drift}} + \underbrace{\sqrt{2T_c} (G^{-1/2})^{jk} dW^j}_{\text{Thermal Noise}}$$

This equation unifies:

1. **Gradient Descent** (on the Value Landscape)
2. **Magnetic Steering** (from Value Curl)
3. **Geodesic Motion** (on the curved Manifold)
4. **Stochastic Exploration** (Langevin Dynamics)

The integration step is modeled as a **Covariant Cross-Attention** layer (Multi-Head Transformer).

#### A. Architecture: Covariant Cross-Attention

- **Mechanism:** Attention heads act as **Wilson Lines** (Parallel Transport operators), comparing queries and keys only after transporting them to a common reference frame (Gauge Invariance).
- **Metric-Temperature:** The softmax temperature is position-dependent:  $\tau(z) \propto 1/\lambda(z)$ . High-curvature regions (large metric  $\lambda$ ) force low temperature (sharp attention), while flat regions allow high temperature (broad exploration).
- **Geodesic Correction:** Christoffel symbols are encoded via linear and quadratic terms in the Query projection.

#### B. Inputs & Outputs (Integration Step):

- **Input:**
  - Current State  $z_t$  (Query position).
  - Action  $u_t$  (Latent Force/Momentum).
  - Memory Context (Keys/Values from past trajectory).
- **Output:**
  - Next State  $z_{t+1}$  (Integrated position after Kick-Drift-Kick).

#### C. Training Losses:

1. **Geodesic Distance Loss:**  $\mathcal{L}_{\text{geo}} \approx (z_{\text{pred}} - z_{\text{true}})^T G(z_t)(z_{\text{pred}} - z_{\text{true}})$ . Minimizes local Riemannian distance. Using the diagonal approximation (Section 5), this becomes a **Weighted MSE**:  $\sum_i G_{ii} (z_{\text{pred}}^{(i)} - z_{\text{true}}^{(i)})^2$ . High-risk dimensions (large  $G_{ii}$ ) are penalized more heavily.
2. **Thermodynamic Consistency:**  $\mathcal{L}_{\text{NLL}} = -\log p(z_{t+1}|z_t, u_t)$ . Ensures the model captures the stochastic thermal noise term ( $\sqrt{2T_c} dW$ ) correctly.

**D. Structural Inductive Bias (Why it acts as an Integrator):** We do not simply train a generic MLP to output  $z_{t+1}$ . Instead, we bake the **Boris-BAOAB** integration scheme directly into the attention mechanism, ensuring the model cannot violate the symplectic structure:

1. **Metric as Temperature:** The attention temperature  $\tau(z) \propto \sqrt{d_k}/\lambda(z)$  forces the update step size to scale inversely with curvature (conformal factor). High-curvature regions (large metric) automatically induce small, cautious steps (sharp attention).
2. **Geodesic Query Terms:** We explicitly feed geometric terms ( $z, z \otimes z$ ) into the Query projection  $Q(z)$ . This forces the attention scores to learn the **Christoffel Symbols**  $\Gamma_{ij}^k$  needed to correct for manifold curvature, rather than making up arbitrary dynamics.
3. **Operator Splitting:** We use **multiple attention heads** to implement the split operators of the BAOAB scheme:
  - *Kick Head (B)*: Updates momentum using force (Gradient + Curl).
  - *Drift Head (A)*: Updates position using momentum (Geodesic flow).
  - *Thermostat Head (O)*: Applies friction and noise (OU process).
  - *Result*: The network is forced to learn a decomposable, reversible integrator rather than a “black box” transition.

## 5. Efficient Metric Computation (Adam-Style)

Computing the full Riemannian metric  $G_{ij}$  ( $D \times D$  tensor) is expensive ( $O(D^3)$  inversion). We use the same engineering tricks as the **Adam Optimizer** to approximate it efficiently ( $O(D)$ ).

### A. The “Adam” Isomorphism:

- **Adam Optimizer:** Maintains a diagonal approximation of the Hessian (via squared gradients) to precondition updates.
  - $v_t = \beta_2 v_{t-1} + (1 - \beta_2) g_t^2$  (Second Moment Estimate).
  - Preconditioner:  $P = \text{diag}(1/\sqrt{v_t})$ .
- **Fragile Agent:** Maintains a diagonal approximation of the Metric Tensor (via Risk Tensor) to curve the space.
  - Metric  $\text{Metric}_t = \beta \text{Metric}_{t-1} + (1 - \beta) \text{Risk}_t$  (Metric Evolution).
  - Geometry:  $G_{ij} \approx \text{diag}(\text{Metric}_t)$ .

### B. Implementation Details:

1. **Diagonal Approximation:** We assume  $G_{ij}$  is diagonal (independent curvature per dimension). This reduces storage from  $O(D^2)$  to  $O(D)$  and inversion to element-wise division.
2. **Risk as Squared Gradient:** The Risk Tensor  $T_{ij}$  is dominated by the gradient of the potential:  $T_{ij} \approx \partial_i \Phi \partial_j \Phi$ . Its diagonal is simply  $(\nabla \Phi)^2$ .
3. **Low-Rank Updates (EMA):** We do not solve the Einstein Field Equations at every step. Instead, we update the metric using an **Exponential Moving Average (EMA)** of the Risk Tensor.
  - *Update Rule:* `metric_diag.lerp_(risk_diag, 1 - momentum)`
  - *Interpretation:* The geometry “flows” slowly towards the high-risk regions, smoothing out transient noise just like Adam smooths gradient variance.

**Result:** We get Riemannian Manifold Hamiltonian Monte Carlo (RMHMC) benefits for the cost of standard SGD+Momentum.