

Single Agent Architecture as Field Theory: Technical TLDR

Guillem Duran Ballester, Jan 2026

1. Latent Space Decomposition

Split-Latent Structure: - Total latent: $Z = (K, z_n, z_{\text{tex}})$ where each component has distinct geometric/physical role - $K \in \{1, \dots, N_c\}$: Discrete macro state (VQ codebook index) - chart assignment on topological atlas - $z_n \in \mathbb{R}^{D_n}$: Continuous nuisance latent - local coordinates within chart (gauge-invariant position) - $z_{\text{tex}} \in \mathbb{R}^{D_t}$: Texture residual - holographic boundary degrees of freedom - Decomposition satisfies $z_e = z_q + z_n + z_{\text{tex}}$ where z_e is raw encoder output, z_q is VQ-quantized macro

Encoder Architecture (AttentiveAtlasEncoder): - Feature extraction: Conv layers ($64 \rightarrow 128 \rightarrow 256$ channels) \rightarrow hidden_dim projection - Cross-attention routing: $w_k = \text{softmax}(\langle K_k, Q(x) \rangle / \sqrt{D})$ where K_k are learnable chart queries - Query projection: $Q(x) = \text{key_proj}(\text{features}(x))$ with LayerNorm stabilization - Chart assignment: $K = \arg \max_k w_k(x)$ - VQ per chart: N_c independent codebooks, each with K_c codes, vectorized quantization - Nuisance extraction: Structure filter $z_n = f_{\text{struct}}(z_e - z_q)$ removes VQ-residual structure - Texture: Holographic residual $z_{\text{tex}} = (z_e - z_q) - z_n$ orthogonal to both macro and nuisance. - **Texture Firewall:** Dynamics depend on z_{macro} and z_n , screening out only z_{tex} .

Training Objectives for Coarse-Graining: To enforce this split-latent structure and ensure valid coarse-graining, we minimize:

$$\mathcal{L}_{\text{latent}} = \mathcal{L}_{\text{VQ}} + \mathcal{L}_{\text{closure}} + \mathcal{L}_{\text{slowness}} + \mathcal{L}_{\text{disentangle}}$$

1. **VQ Loss:** $\|z_e - z_q\|^2 + \beta \|z_q - \text{sg}[z_e]\|^2$. Stabilizes the discrete macro state K .
2. **Causal Enclosure:** $-\log p(K_{t+1}|K_t, a_t)$. Ensures macro dynamics are self-contained (predictable without micro details).
3. **Slowness (Anti-Churn):** $\|e_{K_t} - e_{K_{t-1}}\|_G^2$. Penalizes rapid flickering of the macro state to ensure temporal stability.
4. **Disentanglement:** * *Nuisance KL*: $D_{\text{KL}}(q(z_n|x) \parallel \mathcal{N}(0, I))$. Minimal sufficient nuisance. * *Texture KL*: $D_{\text{KL}}(q(z_{\text{tex}}|x) \parallel \mathcal{N}(0, I))$. Texture should contain no macro info.

Geometric Regularization (Quality Control): To ensure the latent space is well-conditioned (high-fidelity, isometric charts), we add:

$$\mathcal{L}_{\text{reg}} = \mathcal{L}_{\text{VICReg}} + \mathcal{L}_{\text{ortho}}$$

5. **VICReg (Self-Supervision):** Prevents collapse without negative pairs. * *Invariance*: $\|z - z'\|_G^2$. Robustness to view augmentation. * *Variance*: $\max(0, \gamma - \sqrt{\text{Var}(z)})$. Forces code utilization. * *Covariance*: $C(z) \approx I$. Decorrelates latent dimensions (whitening).
6. **Orthogonality (Chart Isometry):** $\|W^T W - I\|$

$I\|_F^2$ on encoder weights. Ensures the mapping from observation to latent space is locally isometric (preserves distances), crucial for meaningful geodesic calculations.

2. The Reward Field & Hodge Decomposition

Physical Context: We model the agent as a particle with **Position and Momentum** performing a **Geodesic Random Walk** on the latent manifold. Because utility is harvested via trajectory traversal, the Reward is naturally defined as a differential **1-form** coupled to velocity, rather than a static scalar field.

Constraint: Reward is not a scalar $r(z)$, but a **1-form** \mathcal{R} that depends on direction ($\mathcal{R}_i \dot{z}^i$). This requires a field-theoretic treatment of value.

Hodge Decomposition: The reward 1-form uniquely decomposes into three orthogonal components:

$$\mathcal{R} = \underbrace{d\Phi}_{\text{Gradient}} + \underbrace{\delta\Psi}_{\text{Solenoidal}} + \underbrace{\eta}_{\text{Harmonic}}$$

- **Gradient** (Φ): Optimizable scalar potential (Conservative Value).
- **Solenoidal** (Ψ): Vector potential generating the **Value Curl** (Magnetic Field) $\mathcal{F} = d\mathcal{R} = d\delta\Psi$. - *Physics*: Just as a magnetic field B exerts a **Lorentz Force** $v \times B$, the Value Curl \mathcal{F} exerts a velocity-dependent force $\mathcal{F}_{ij} \dot{z}^j$ that steers the agent to **orbit** value cycles rather than converge.
- **Harmonic** (η): Topological cycles (manifold holes).

The Screened Poisson Equation (Conservative Case): When the **Value Curl** vanishes ($\mathcal{F} = 0$), the scalar value function $V(z)$ satisfies the **Helmholtz Equation** on the manifold:

$$(-\Delta_G + \kappa^2)V(z) = \rho_r(z)$$

where Δ_G is the Laplace-Beltrami operator, κ is the screening mass, and ρ_r is the reward source density.

The Composite Navigation Potential (Runtime): The agent navigates a **Composite Potential** Φ_{eff} constructed at runtime by summing learned and intrinsic signals:

$$\Phi_{\text{eff}}(z) = \underbrace{V(z)}_{\text{Learned}} + \underbrace{U(z)}_{\text{Intrinsic}} + \underbrace{\mathcal{B}_{\text{safety}}(z)}_{\text{Fixed}}$$

1. **Control** (V): **Learned** scalar value. Drives the agent towards high-reward regions. * *Loss*: Helmholtz Residual (see below).
2. **Generation** (U): **Intrinsic** entropy potential (e.g., Hyperbolic expansion $-\log \text{vol}(z)$). Drives exploration away from the origin. * *Loss*: None (Fixed geometric prior).
3. **Safety** ($\mathcal{B}_{\text{safety}}$): **Fixed** safety barrier. Hard constraints (e.g., capacity limits) modeled as high-energy walls. * *Loss*: None (Fixed constraint).

Neural Hodge Decomposition (Implementation): We approximate the Hodge components using a **Multi-Head Network** sharing a common feature backbone:

1. **Scalar Head (Φ):** Outputs scalar $V(z)$.
 - *Loss:* Helmholtz Residual on the symmetric part of the reward.
 - $\mathcal{L}_\Phi = \|V(z) - (r_{\text{sym}} + \gamma V(z'))\|^2$.
2. **Solenoidal Head (Ψ):** Outputs vector potential $A(z) \in \mathbb{R}^D$ (where $\mathcal{F} = dA$).
 - *Loss:* Residual reconstruction on the antisymmetric part.
 - $\mathcal{L}_\Psi = \|\langle \mathcal{R}, v \rangle - (\langle \nabla \Phi, v \rangle + \langle \nabla \times A, v \rangle)\|^2$. The curl absorbs the non-integrable reward residual.
3. **Harmonic Head (η):** A set of **learnable constant 1-forms** η_k per chart k .
 - *Mechanism:* Captures global topological currents (net flux through manifold holes) that are locally constant.
 - *Loss:* Projected residual after removing Gradient and Curl components.

Value Function Objectives: To define the value field $V(z)$ as the solution to the Screened Poisson Equation, we minimize:

$$\mathcal{L}_{\text{critic}} = \underbrace{\|V(z) - (r + \gamma V(z'))\|^2}_{\text{Helmholtz Residual (TD)}} + \lambda_{\text{geo}} \underbrace{\|\nabla_G V\|^2}_{\text{Smoothness}}$$

1. **Helmholtz Residual:** Enforces the PDE source term (approx. Bellman error).
2. **Geometric Regularization:** Enforces field smoothness with respect to the manifold metric ($\|\nabla_G V\|^2 = G^{ij} \partial_i V \partial_j V$).

3. The Policy Network (Latent Action)

The Policy acts as an **External Force Field** $u_\pi(z)$ pushing the agent through the latent manifold. It operates in a **Latent Action Space** (the Tangent Bundle $T\mathcal{Z}$), decoupling low-level motor commands from high-level intent.

A. The Policy Model (π_ϕ): - **Role:** Symmetry-breaking control field. Converts potential energy into kinetic motion. - **Input:** Latent State $z_t \in \mathcal{Z}$ (Position). - **Output:** Latent Force/Action $u_t \in T_{z_t} \mathcal{Z}$ (Tangent Vector). - **Note:** This latent force is subsequently decoded into boundary motor torques action a_t by the Motor/Action Decoder (Neumann Condition). - **Latent Action Space:** The Tangent Bundle $T\mathcal{Z}$. Actions are vectors “pushing” the state along geodesics.

B. Training Losses (Tier 1): 1. **Task Loss ($\mathcal{L}_{\text{task}}$):** Standard Policy Gradient / Reinforce objective to maximize expected returns. 2. **Entropy Bonus (Ergodicity):** $\mathcal{L}_{\text{ent}} = -H(\pi)$. Penalizes low entropy distributions to prevent premature mode collapse and ensure thermodynamic equilibrium. 3. **Zeno Penalty (Temporal Smoothness):** $\mathcal{L}_{\text{zeno}} = D_{\text{KL}}(\pi_t \parallel \pi_{t-1})$. Penalizes

infinite-frequency oscillations (Zeno behavior) to ensure physically realizable trajectories.

4. The World Model (Covariant Integrator)

We define the World Model not as a generic RNN, but as a **Neural Integrator** that approximates the **Lorentz-Langevin SDE**:

$$dz^k = \underbrace{(-G^{kj}\partial_j\Phi + u_\pi^k)}_{\text{Gradient} + \text{Policy}} ds + \underbrace{\beta G^{km}\mathcal{F}_{mj}\dot{z}^j ds}_{\text{Lorentz Force}} - \underbrace{\Gamma_{ij}^k z^i \dot{z}^j ds}_{\text{Geodesic Drift}} + \underbrace{\sqrt{2T_c}(G^{-1/2})^{jk} dW^j}_{\text{Thermal Noise}}$$

This equation unifies: 1. **Gradient Descent** (on the Value Landscape) 2. **Magnetic Steering** (from Value Curl) 3. **Geodesic Motion** (on the curved Manifold) 4. **Stochastic Exploration** (Langevin Dynamics)

The integration step is modeled as a **Covariant Cross-Attention** layer (Multi-Head Transformer).

A. Architecture: Covariant Cross-Attention - Mechanism: Attention heads act as **Wilson Lines** (Parallel Transport operators), comparing queries and keys only after transporting them to a common reference frame (Gauge Invariance). - **Metric-Temperature:** The softmax temperature is position-dependent: $\tau(z) \propto 1/\lambda(z)$. High-curvature regions (large metric λ) force low temperature (sharp attention), while flat regions allow high temperature (broad exploration). - **Geodesic Correction:** Christoffel symbols are encoded via linear and quadratic terms in the Query projection.

B. Inputs & Outputs (Integration Step): - **Input:** - Current State z_t (Query position). - Action u_t (Latent Force/Momentum). - Memory Context (Keys/Values from past trajectory). - **Output:** - Next State z_{t+1} (Integrated position after Kick-Drift-Kick).

C. Training Losses: 1. **Geodesic Distance Loss:** $\mathcal{L}_{\text{geo}} \approx (z_{\text{pred}} - z_{\text{true}})^T G(z_t)(z_{\text{pred}} - z_{\text{true}})$. Minimizes local Riemannian distance. Using the diagonal approximation (Section 5), this becomes a **Weighted MSE**: $\sum_i G_{ii}(z_{\text{pred}}^{(i)} - z_{\text{true}}^{(i)})^2$. High-risk dimensions (large G_{ii}) are penalized more heavily. 2. **Thermodynamic Consistency:** $\mathcal{L}_{\text{NLL}} = -\log p(z_{t+1}|z_t, u_t)$. Ensures the model captures the stochastic thermal noise term ($\sqrt{2T_c}dW$) correctly.

D. Structural Inductive Bias (Why it acts as an Integrator): We do not simply train a generic MLP to output z_{t+1} . Instead, we bake the **Boris-BAOAB** integration scheme directly into the attention mechanism, ensuring the model cannot violate the symplectic structure:

1. **Metric as Temperature:** The attention temperature $\tau(z) \propto \sqrt{d_k}/\lambda(z)$ forces the update step size to scale inversely with curvature (conformal

factor). High-curvature regions (large metric) automatically induce small, cautious steps (sharp attention).

2. **Geodesic Query Terms:** We explicitly feed geometric terms ($z, z \otimes z$) into the Query projection $Q(z)$. This forces the attention scores to learn the **Christoffel Symbols** Γ_{ij}^k needed to correct for manifold curvature, rather than making up arbitrary dynamics.
3. **Operator Splitting:** We use **multiple attention heads** to implement the split operators of the BAOAB scheme:
 - *Kick Head (B):* Updates momentum using force (Gradient + Curl).
 - *Drift Head (A):* Updates position using momentum (Geodesic flow).
 - *Thermostat Head (O):* Applies friction and noise (OU process).
 - *Result:* The network is forced to learn a decomposable, reversible integrator rather than a “black box” transition.

5. Efficient Metric Computation (Adam-Style)

Computing the full Riemannian metric G_{ij} ($D \times D$ tensor) is expensive ($O(D^3)$ inversion). We use the same engineering tricks as the **Adam Optimizer** to approximate it efficiently ($O(D)$).

A. The “Adam” Isomorphism: - Adam Optimizer: Maintains a diagonal approximation of the Hessian (via squared gradients) to precondition updates.
- $v_t = \beta_2 v_{t-1} + (1 - \beta_2)g_t^2$ (Second Moment Estimate). - Preconditioner: $P = \text{diag}(1/\sqrt{v_t})$. - **Fragile Agent:** Maintains a diagonal approximation of the Metric Tensor (via Risk Tensor) to curve the space. - Metric $_t = \beta \text{Metric}_{t-1} + (1 - \beta) \text{Risk}_t$ (Metric Evolution). - Geometry: $G_{ij} \approx \text{diag}(\text{Metric}_t)$.

B. Implementation Details: 1. **Diagonal Approximation:** We assume G_{ij} is diagonal (independent curvature per dimension). This reduces storage from $O(D^2)$ to $O(D)$ and inversion to element-wise division. 2. **Risk as Squared Gradient:** The Risk Tensor T_{ij} is dominated by the gradient of the potential: $T_{ij} \approx \partial_i \Phi \partial_j \Phi$. Its diagonal is simply $(\nabla \Phi)^2$. 3. **Low-Rank Updates (EMA):** We do not solve the Einstein Field Equations at every step. Instead, we update the metric using an **Exponential Moving Average (EMA)** of the Risk Tensor. - *Update Rule:* `metric_diag.lerp_(risk_diag, 1 - momentum)`
- *Interpretation:* The geometry “flows” slowly towards the high-risk regions, smoothing out transient noise just like Adam smooths gradient variance.

Result: We get Riemannian Manifold Hamiltonian Monte Carlo (RMHMC) benefits for the cost of standard SGD+Momentum.